

# TORCS Gaming AI Agent: Q-Learning-Based Autonomous Racecar Control

Kalbe Raza (i220794-B)  
Abdullah Mehmood (i220978-B)  
Aneeq Malik (i221167-B)  
Artificial Intelligence Course Project

## Abstract

This report presents the "TORCS Gaming AI Agent," a Q-learning-based system developed for an Artificial Intelligence course to control an autonomous racecar in The Open Racing Car Simulator (TORCS). The system comprises two components: **MultiTrackQLearningTrainer**, which trains a Q-table offline using track data from CSV files, and **QLearningDriver**, which deploys the Q-table in real-time to navigate TORCS tracks. The system optimizes racing performance by balancing speed, track position, and safety, using Q-learning for decision-making and rule-based logic for edge cases like getting stuck or excessive swerving. Detailed explanations of each function, their integration, and the rationale for choosing Q-learning over neural networks are provided, along with a graph illustrating training performance.

## 1 Overview

This report details the "TORCS Gaming AI Agent," a project developed for an Artificial Intelligence course, focusing on autonomous racecar control in TORCS using Q-learning. The system consists of two components:

- **MultiTrackQLearningTrainer:** Trains a Q-table offline using track data from CSV files (`GL.csv`, `DL.csv`, `TL.csv`) to learn optimal driving actions across multiple TORCS tracks.
- **QLearningDriver:** Deploys the trained Q-table in real-time to control a racecar in TORCS, using rule-based control as a fallback for robustness.

The system optimizes racing performance by balancing speed, track position, and safety, leveraging Q-learning for decision-making and rule-based logic for edge cases like getting stuck or excessive swerving. The report includes detailed explanations of each function in both classes, their integration, and the rationale for choosing Q-learning over neural networks for this AI-driven project.

## 2 MultiTrackQLearningTrainer: Function Explanations

The `MultiTrackQLearningTrainer` class trains a Q-table using Q-learning on offline track data from TORCS, employing multithreading for efficiency. Below is a detailed explanation of each function, with naive analogies for clarity.

### 2.1 `__init__(self, track_files=None, q_table_file='multi_track_q_table.pkl', batch_size=20000)`

**Purpose:** Initializes the trainer with parameters and data structures.

**Detailed Explanation:**

- **Parameters:**

- `track_files`: List of CSV files (defaults to `['GL1.csv', 'DL1.csv', 'TL1.csv']`).
- `q_table_file`: File for saving/loading the Q-table (`'multi_track_q_table.pkl'`).
- `batch_size`: Number of data points per episode (20,000).

- **Attributes:**

- `q_table`: `defaultdict` with 9 actions (3 steering  $\times$  3 accel/brake).
- `alpha` (0.1): Learning rate for Q-table updates.
- `gamma` (0.95): Discount factor for future rewards.
- `epsilon` (0.8): Exploration rate, decaying to 0.01 with 0.995 decay.
- `state_bins`: Bins for discretizing state variables.
- `actions`: 9 action combinations (steering:  $[-0.5, 0.0, 0.5]$ , accel/brake:  $[(1.0, 0.0), (0.5, 0.0), (0.0, 0.0), (0.0, 0.5), (0.0, 1.0), (-0.5, 0.0), (-0.5, 0.5), (-0.5, 1.0)]$ ).
- `q_table_lock`: Thread-safe lock for Q-table updates.
- `state_visit_counts`: Tracks state visits for adaptive learning rate.
- `track_data`: Stores loaded track DataFrames.

- **Functionality:** Initializes bins, actions, and loads the Q-table if available.

**Naive Analogy:** A kid learning to race toy cars in TORCS gets a notebook (`q_table`) to record good moves. They're told to learn slowly (`alpha`), think about future laps (`gamma`), and try new moves often (`epsilon`). The track is divided into zones (`state_bins`), and they have 9 ways to drive (`actions`). A lock ensures only one kid writes in the notebook at a time.

### 2.2 `define_state_bins(self)`

**Purpose:** Defines bins to discretize continuous state variables.

- Returns a dictionary with bins for:

- `SpeedX`:  $-50$  to  $300$  (25 bins,  $\sim 14.2$  units).
- `TrackPosition`:  $-2.0$  to  $2.0$  (10 bins,  $\sim 0.4$  units).

- **Angle:**  $-\pi$  to  $\pi$  (12 bins,  $\sim 0.52$  radians).
- **Track<sub>g</sub> :** 0 to 250 (8 bins,  $\sim 31.25$  units).
- **Damage:** 0 to 10,000 (6 bins,  $\sim 1666.67$  units).

- Uses `np.linspace` for evenly spaced bins.

**Naive Analogy:** The kid divides the TORCS track into a grid, like a board game, grouping speeds into “slow,” “medium,” “fast,” etc., so they can easily note what to do in each grid square.

## 2.3 `define_actions(self)`

**Purpose:** Defines the action space.

**Detailed Explanation:**

- Creates 9 actions combining 3 steering options ( $-0.5, 0.0, 0.5$ ) with 3 accel/brake options (full, half, none).

**Naive Analogy:** The kid’s toy car in TORCS has buttons for turning left, straight, or right, and pushing hard, lightly, or not at all, giving 9 possible moves.

## 2.4 `discretize_state(self, state)`

**Purpose:** Converts continuous state to discrete state tuple.

**Detailed Explanation:**

- Maps each state variable to a bin index using `np.digitize`, ensuring indices stay within bounds.
- Returns a tuple (e.g., (8, 4, 6, 2, 0)).

**Naive Analogy:** The kid looks at their car’s speed and position in TORCS, then labels it as “fast, near center, slight left turn” to check their notebook.

## 2.5 `compute_reward(self, state, nextstate)`

**Purpose:** Calculates reward for state transitions.

**Detailed Explanation:**

- **Rewards:** Speed ( $+0.1/\text{unit}$ ), acceleration ( $+5.0$ ), centering ( $+2.0$ ), progress ( $+10.0/\text{unit}$ ).
- **Penalties:** Damage ( $-10.0/\text{unit}$ ), off-track ( $-5.0$  if  $|\text{TrackPos}| > 1.0$ ,  $-15.0$  if  $> 1.5$ ), backward speed ( $-20.0$ ), angle ( $-2.0/\text{radian}$ ).

**Naive Analogy:** The kid in TORCS gets points for going fast, speeding up, staying on the path, and moving forward, but loses points for crashing, swerving, or going backward.

## 2.6 `chooseaction(self, state)`

**Purpose:** Selects an action using epsilon-greedy policy.

**Detailed Explanation:**

- With probability `epsilon` (0.8), picks a random action; otherwise, chooses the action with the highest Q-value.

**Naive Analogy:** The kid flips a coin in TORCS: 80% chance they try a random move, 20% chance they pick the best move from their notebook.

## 2.7 `updateqtable(self, state, action, reward, nextstate)`

**Purpose:** Updates the Q-table using Q-learning.

**Detailed Explanation:**

- Discretizes states, uses thread-safe lock, adjusts learning rate based on state visits, and applies:

$$Q(s, a) = Q(s, a) + \alpha \cdot (\text{reward} + \gamma \cdot \max_{a'}(Q(s', a')) - Q(s, a))$$

**Naive Analogy:** The kid in TORCS updates their notebook after a move, blending new points (`reward`) with future possibilities (`max(Q(s', a'))`), writing carefully to avoid mistakes.

## 2.8 `loadqtable(self)`

**Purpose:** Loads a saved Q-table.

**Detailed Explanation:**

- Loads `qtablefileintoadefaultdict, logging the number of states`. **Naive Analogy:** The kid loads their old TORCS notebook to recall past lessons.

## 2.9 `saveqtable(self)`

**Purpose:** Saves the Q-table to a file.

**Detailed Explanation:**

- Converts `qtabletoadictionary, saves it with pickle, using a thread-safe lock`. **Naive Analogy:** The kid saves their TORCS notebook for future races, ensuring no one else writes in it during copying.

## 2.10 `analyzestates(self, df, trackname = None)`

**Purpose:** Analyzes state distribution in track data.

**Detailed Explanation:**

- Counts occurrences in `SpeedX` and `TrackPosition` bins to log coverage (e.g., how many high-speed samples).

**Naive Analogy:** The kid checks how often they drove fast or near the edge in TORCS to ensure they practiced all scenarios.

### 2.11 `load_track_thread(self, track_file, results_queue)`

**Purpose:** Loads a track's CSV data in a thread.

**Detailed Explanation:**

- Reads CSV, analyzes state distribution, and puts the DataFrame in `results_queue`. **Naive Analogy:** A helper reads a TORCS track map, checks its details, and puts it in a shared box for the kid.

### 2.12 `load_track_data(self)`

**Purpose:** Coordinates parallel loading of track data.

**Detailed Explanation:**

- Starts threads for each track, collects DataFrames from `results_queue`, and stores them in `Naive Analogy:` Multiple helpers grab TORCS track maps simultaneously, organizing them in a folder for the kid.

### 2.13 `train_on_track(self, track_name, episode, episodes)`

**Purpose:** Trains on a track for one episode.

**Detailed Explanation:**

- Samples `batch_size`, `rows`, `add_high-speed samples`, `processes state pairs to choose actions`, `compute reward table`. **Naive Analogy:** The kid practices one lap in TORCS, trying moves, earning points, and updating their notebook, focusing extra on fast moments.

### 2.14 `evaluate_on_track(self, track_name, episode)`

**Purpose:** Evaluates performance on a track (not used in training).

**Detailed Explanation:**

- Tests on validation data, computing average reward without updating the Q-table.

**Naive Analogy:** The kid tests their skills in TORCS by driving without writing in the notebook, checking how many points they get.

### 2.15 `train_track_thread(self, track_name, episode, episodes, results_queue)`

**Purpose:** Trains on a track in a thread.

**Detailed Explanation:**

- Calls `train_on_track`, `put total reward in` **Naive Analogy:** A helper teaches the kid on one TORCS track and reports their score to a shared box.

### 2.16 `train(self, episodes=100, max_workers=3)`

**Purpose:** Orchestrates training across tracks and episodes.

**Detailed Explanation:**

- Loads track data, trains in parallel (max 3 threads), decays `epsilon`, saves Q-table every 20 episodes, and logs statistics.

**Naive Analogy:** The kid practices on all TORCS tracks for 100 days, with helpers teaching in parallel, saving their notebook regularly.

## 3 QLearningDriver: Function Explanations

The `QLearningDriver` class controls a racecar in real-time within TORCS using the trained Q-table, with rule-based control for unexplored states. Below are the function explanations.

### 3.1 `__init__(self, stage)`

**Purpose:** Initializes the driver for real-time control in TORCS.

**Detailed Explanation:**

- Sets up Q-table (15 actions), finer bins (20 for `Angle`), `epsilon=0.1`, and parameters for steering, gear, stuck recovery, and oscillation detection.

**Naive Analogy:** The kid prepares to race in TORCS with their notebook, a better controller (15 buttons), and tools to avoid getting stuck or wobbling.

### 3.2 `init(self)`

**Purpose:** Signals readiness to the TORCS simulator.

**Detailed Explanation:** Returns "`init racer`".

**Naive Analogy:** The kid says, "Ready to race!" to start the TORCS game.

### 3.3 `define_state_bins(self)`

**Purpose:** Defines bins for state discretization in TORCS.

**Detailed Explanation:**

- Similar to trainer, but with 20 bins for `Angle` for precision.

**Naive Analogy:** The kid uses a finer grid for turns in TORCS to make smarter steering choices.

### 3.4 `define_actions(self)`

**Purpose:** Defines 15 actions for TORCS control.

**Detailed Explanation:**

- Combines 5 steering options ( $-0.5, -0.25, 0.0, 0.25, 0.5$ ) with 3 accel/brake options.

**Naive Analogy:** The kid's TORCS controller has more turn options for smoother driving.

### 3.5 `load_qtable(self)`

**Purpose:** Loads the Q-table for TORCS.

**Detailed Explanation:**

- Loads `qtable.pkl`, `logsstatesandsampleQ-values`. **Naive Analogy:** The kid opens their TORCS notebook to use practice lessons.

### 3.6 onShutDown(self)

**Purpose:** Logs state visit statistics on TORCS shutdown.

**Detailed Explanation:** Logs top 10 most visited states.

**Naive Analogy:** At the end of a TORCS race, the kid checks which track spots they visited most.

### 3.7 onRestart(self)

**Purpose:** Resets state for a new TORCS race.

**Detailed Explanation:** Clears previous state, steering, gear, and counters.

**Naive Analogy:** The kid starts a new TORCS race, forgetting the last one's details.

### 3.8 get<sub>s</sub>tate<sub>from\_s</sub>string(self, string)

**Purpose:** Parses TORCS sensor data into a state dictionary.

**Detailed Explanation:**

- Extracts Angle, SpeedX, TrackPosition, Track<sub>0</sub>, Damage, Gear, RPM, DistanceCovered, TrackFront, Tr

**Naive Analogy:** The kid reads their TORCS car's dashboard to know speed, position, and obstacles.

### 3.9 discretize<sub>s</sub>tate(self, state)

**Purpose:** Converts TORCS state to discrete tuple.

**Detailed Explanation:** Same as trainer, mapping to bin indices.

**Naive Analogy:** The kid labels their situation in TORCS (e.g., "fast, centered, slight turn") to check their notebook.

### 3.10 detect<sub>s</sub>tuck(self, state)

**Purpose:** Detects if the car is stuck in TORCS.

**Detailed Explanation:**

- Increments stuck<sub>counter</sub> if SpeedX < 3.0; enters **Naive Analogy:** If the car barely moves in TORCS, the kid knows it's stuck and plans to wiggle free.

### 3.11 recover<sub>from\_s</sub>tuck(self)

**Purpose:** Executes recovery actions in TORCS.

**Detailed Explanation:**

- Applies steering ( $\pm 0.3$ ) and acceleration (1.0) for 30 steps.

**Naive Analogy:** The kid steers and pushes the pedal in TORCS to get out of a ditch.

### 3.12 detect<sub>o</sub>scillation(self, trackPos)

**Purpose:** Detects excessive swerving in TORCS.

**Detailed Explanation:**

- Tracks TrackPosition, counts center crossings; increases damping if  $\geq 4$  in 6 steps.

**Naive Analogy:** If the kid zigzags too much in TORCS, they drive smoother to stabilize.

### 3.13 `apply_safety_checks(self, steer, accel, brake, track_sensors, speed)`

**Purpose:** Adjusts actions to avoid collisions in TORCS.

**Detailed Explanation:**

- Brakes or reduces acceleration if front sensors show obstacles (15 or 7 units).

**Naive Analogy:** The kid slows down in TORCS if a wall is close to avoid crashing.

### 3.14 `rule_based_control(self, state)`

**Purpose:** Provides fallback control in TORCS.

**Detailed Explanation:**

- Computes steering (angle + track position + curve adjustments, damped), speed (targets: 200, 140, 90), and gear (RPM-based).

**Naive Analogy:** If the notebook's empty, the kid drives by instinct in TORCS, steering to stay centered and slowing for turns.

### 3.15 `drive(self, string)`

**Purpose:** Controls the car in real-time in TORCS.

**Detailed Explanation:**

- Parses state, checks for stuck conditions, uses Q-table (if  $\max(Q) > 0.3$ ) or rule-based control, applies safety checks, and outputs control string.

**Naive Analogy:** The kid reads the dashboard in TORCS, checks their notebook or uses instinct, avoids crashes, and tells the car what to do.

## 4 Integration of Trainer and Driver

### 4.1 How They Work Together

- **Training (MultiTrackQLearningTrainer):**
  - Processes offline CSV data from multiple TORCS tracks.
  - Discretizes states, defines 9 actions, and trains a Q-table using Q-learning.
  - Saves Q-table to `multi_track_qtable.pkl`.
- **Deployment (QLearningDriver):**
  - Loads Q-table (from `qtable.pkl`, requiring filename alignment). Uses finer bins (20 for Angle).
  - Parses real-time TORCS sensor data, applies Q-table actions or rule-based control, and ensures safety.
- **Integration Steps:**
  - Align Q-table file names (e.g., rename to `qtable.pkl`). Update trainer's actionspace to 15 actions.
  - The trainer's Q-table provides generalized knowledge, enabling the driver to adapt to similar TORCS tracks in real-time.



## 4.2 Challenges and Solutions

- **Action Space Mismatch:** Trainer (9 actions) vs. driver (15 actions). *Solution:* Update trainer to use driver's actions.
- **File Name Mismatch:** Trainer saves to `multitrackqtable.pkl`, driver loads from

## 5 Training Performance Graph

Figure 1 illustrates the average reward per episode during training across multiple TORCS tracks, as processed by `MultiTrackQLearningTrainer`. The graph shows the learning progress, with rewards increasing as the Q-table converges, reflecting improved driving performance in the TORCS environment.




Figure 1: Average reward per episode during Q-learning training across multiple TORCS tracks, demonstrating convergence of the Q-table.

*Note:* The figure is a placeholder for a graph of training rewards. Replace `training_rewards.png` with

## 6 Why Q-Learning Instead of Neural Networks?

### 6.1 Advantages of Q-Learning

- **Simplicity and Interpretability:** Q-learning uses a table, making it easy to inspect Q-values (logged in `loadqtable`). *Neural networks are complex and opaque.* Manage *Discretized states* ( $\sim 25 \times 10 \times 20 \times 8 \times 6$ ) and 15 actions fit in memory, enabling fast lookups. Neural networks suit high-dimensional inputs (e.g., images).
- **Data Efficiency:** Q-learning trains effectively with limited CSV data (20,000 rows/episode). Neural networks require large datasets and extensive training.
- **Low Computational Cost:** Q-learning runs on modest hardware, with multithreading in the trainer and lightweight lookups in the driver. Neural networks need GPUs.

- Robust Fallback: The driver's rule-based control handles unexplored states, stuck situations, and oscillation in TORCS, ensuring reliability. Neural networks need complex architectures (e.g., DQN) for similar robustness.
- Exploration Control: Epsilon-greedy ( $\epsilon = 0.8$  to  $0.01$  in trainer,  $0.1$  in driver) is simple to tune. Neural networks require advanced exploration strategies.

## 6.2 Conclusion

Q-learning is chosen for its simplicity, efficiency, and suitability for a discretized state-action space in the TORCS environment. The hybrid approach (Q-table + rule-based control) ensures robust real-time performance, making it ideal for this AI course project.

## 7 Conclusion

The MultiTrackQLearningTrainer and QLearningDriver form a cohesive system for the "TORCS Gaming AI Agent," developed as an Artificial Intelligence course project. The trainer builds a robust Q-table offline, while the driver applies it in real-time within TORCS, augmented by rule-based logic for safety and adaptability. Q-learning's simplicity and efficiency make it a fitting choice over neural networks for this application, given the manageable state space and data constraints.